

Synchronous Shared Memory Parallel Examples

HPC Spring 2017

Prof. Robert van Engelen





Examples

- Data parallel prefix sum and OpenMP example
- Task parallel prefix sum and OpenMP example
- Simple heat distribution problem with OpenMP
- Iterative solver with OpenMP
- Simple heat distribution problem with HPF
- Iterative solver with HPF
- Gaussian Elimination with HPF



Dataparallel Prefix Sum

Serial

```
sum[0] = x[0];  
for (i = 1; i < n; i++)  
    sum[i] = sum[i-1] + x[i];
```

Dataparallel

```
for (j = 0; j < log2(n); j++)  
    forall (i = 2j; i < n; i++)  
        x[i] = x[i] + x[i - 2j];
```

***Dataparallel =
synchronous
(lock-step)***

Time complexity?

***Parallel efficiency
and speedup?***

*Dataparallel forall: concurrent write and read, but read always fetches old value:
forall has “copy-in-copy-out” semantics (viz. CRCW PRAM model)*

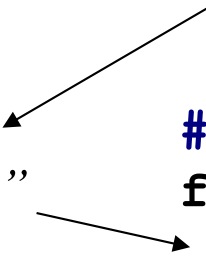


OpenMP Prefix Sum v1

```
for (j = 0; j < log2(n); j++)
{
    #pragma omp parallel private(i)
    {
        #pragma omp for
        for (i = 1<<j; i < n; i++)
            t[i] = x[i] + x[i - 1<<j];

        #pragma omp for
        for (i = 1<<j; i < n; i++)
            x[i] = t[i];
    }
}
```

“Copy out”



*What about
overhead:
when/where
would it be
better to
create/join
threads?*

Note: use bitshift to compute $2^j = 1 \ll j$



Task Parallel Prefix Sum

```
for each processor  $0 \leq p < n$ 
private j
{
    for (j = 1; j < n; j = 2*j)
    {
        if (p >= j)
            t[p] = x[p] + x[p-j];
        barrier;
        x[p] = t[p];
        barrier;
    }
}
```

Task/thread-parallel: parallelize outer loops, each processor takes an outer loop iteration to execute concurrently



OpenMP Prefix Sum v2

```
#pragma omp parallel shared(n,x,t) private(j,tid) num_threads(n)
{
    tid = omp_get_thread_num();

    for (j = 1; j < n; j = 2*j)
    {
        if (tid >= j)
            t[tid] = x[tid] + x[tid - j];
        #pragma omp barrier

        x[tid] = t[tid];
        #pragma omp barrier
    }
}
```

*Alternative to **private(j,tid)**:
declare **j** and **tid** locally in this block*

*Uses n threads!
What if n is really large?*



OpenMP Prefix Sum v3

```
#pragma omp parallel shared(n,nthr,x,z) private(i,j,tid,work,lo,hi)
{
    #pragma omp single
        nthr = omp_get_num_threads();
        tid = omp_get_thread_num();
        work = (n + nthr-1) / nthr;
        lo = work * tid;
        hi = lo + work;
        if (hi > n)
            hi = n;
    for (i = lo+1; i < hi; i++)
        x[i] = x[i] + x[i-1];
    z[tid] = x[hi-1];
    #pragma omp barrier
    for (j = 1; j < nthr; j = 2*j)
    {
        if (tid >= j)
            z[tid] = z[tid] + z[tid - j];
        #pragma omp barrier
    }
    for (i = lo; i < hi; i++)
        x[i] = x[i] + z[tid] - x[hi-1];
}
```

Note: assumes nthreads = 2^k

Local prefix sum over x

$z = \text{local prefix sum } x[\text{hi}]$

Global prefix sum over z

Update local prefix sum x



Dataparallel Heat Distribution Problem

$$h_{i,j} = \frac{h_{i-1,j} + h_{i+1,j} + h_{i,j-1} + h_{i,j+1}}{4}$$

```
for (iter = 0; iter < limit; iter++)  
  forall (i = 0; i < n; i++)  
    forall (j = 0; j < n; j++)  
      h[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
```

Dataparallel = synchronous

Corresponds to Jacobi iteration

*Dataparallel forall: concurrent write and read, but read always fetches old value:
forall has “copy-in-copy-out” semantics (viz. CRCW PRAM model)*



OpenMP Heat Distribution Problem

```
#pragma omp parallel shared(newh,h,limit,n) private(iter,i,j)
{
    for (iter = 0; iter < limit; iter++)
    {
        #pragma omp for
        for (i = 0; i < n; i++)
            for (j = 0; j < n; j++)
                newh[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
        #pragma omp for
        for (i = 0; i < n; i++)
            for (j = 0; j < n; j++)
                h[i][j] = newh[i][j];
    }
}
```

Corresponds to Jacobi iteration



Dataparallel Heat Distribution

Red-black Ordering

```
for (iter = 0; iter < limit; iter++)
{
    forall (i = 0; i < n; i++)
        forall (j = 0; j < n; j++)
            if ((i+j) % 2 != 0)
                h[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
    forall (i = 0; i < n; i++)
        forall (j = 0; j < n; j++)
            if ((i+j) % 2 == 0)
                h[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
}
```

Dataparallel = synchronous



OpenMP Heat Distribution Red-black Ordering

```
#pragma omp parallel shared(h,limit,n) private(iter,i,j)
{
    for (iter = 0; iter < limit; iter++)
    {
        #pragma omp for
        for (i = 0; i < n; i++)
            for (j = 0; j < n; j++)
                if ((i+j) % 2 != 0)
                    h[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
        #pragma omp for
        for (i = 0; i < n; i++)
            for (j = 0; j < n; j++)
                if ((i+j) % 2 == 0)
                    h[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
    }
}
```



Iterative Solver

$$\mathbf{Ax} = \mathbf{b}$$

Jacobi iteration

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

Stopping criteria:

$$\sqrt{\sum_{i=0}^{n-1} (x_i^k - x_i^{k-1})^2} < \epsilon$$

Pacheco

$$\left| \sum_{j=0}^{n-1} a_{i,j} x_j^k - b_i \right| < \epsilon \quad \forall i = 0, \dots, n-1$$

Bertsekas and Tsitsiklis



Iterative Solver: Jacobi Method

```
for (i = 0; i < n; i++)  
    x[i] = b[i];
```

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

```
for (iter = 0; iter < limit; iter++)  
{  
    for (i = 0; i < n; i++)  
    {  
        sum = -a[i][i] * x[i]; // correction to sum over j!=i  
        for (j = 0; j < n; j++)  
            sum = sum + a[i][j] * x[j];  
        new_x[i] = (b[i] - sum) / a[i][i];  
    }  
    for (i = 0; i < n; i++)  
        x[i] = new_x[i];  
}
```

Note: stopping criterium omitted



Iterative Solver: Jacobi Method

```
for (i = 0; i < n; i++)  
    x[i] = b[i];
```

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

```
for (iter = 0; iter < limit; iter++)  
{  
    for (i = 0; i < n; i++)  
        sum[i] = -a[i][i] * x[i]; // correction to sum over j!=i  
    for (i = 0; i < n; i++)  
        for (j = 0; j < n; j++)  
            sum[i] = sum[i] + a[i][j] * x[j];  
    for (i = 0; i < n; i++)  
        new_x[i] = (b[i] - sum[i]) / a[i][i];  
    for (i = 0; i < n; i++)  
        x[i] = new_x[i];  
}
```

Note: after array expansion of scalar sum and loop fission

Note: stopping criterium omitted



Dataparallel Iterative Solver: Jacobi Method

```
for (i = 0; i < n; i++)  
    x[i] = b[i];
```

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

```
for (iter = 0; iter < limit; iter++)  
{  
    forall (i = 0; i < n; i++)  
        sum[i] = -a[i][i] * x[i];  
    for (j = 0; j < n; j++)  
        forall (i = 0; i < n; i++)  
            sum[i] = sum[i] + a[i][j] * x[j];  
    forall (i = 0; i < n; i++)  
        x[i] = (b[i] - sum[i]) / a[i][i];  
}
```

*Note: after loop interchange and
forall-parallelization*

Dataparallel = synchronous (lock-step)

Note: stopping criterium omitted



Task Parallel Iterative Solver: Jacobi Method

for each processor $0 \leq p < n$

private iter, sum, j

{

 x[p] = b[p];

 for (iter = 0; iter < limit; iter++)

 {

 sum = -a[p][p] * x[p];

 for (j = 0; j < n; j++)

 sum = sum + a[p][j] * x[j];

barrier;

 x[p] = (b[p] - sum) / a[p][p];

barrier;

 }

}

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

*Note: each processor is assigned
to an iteration i*

Note: stopping criterium omitted



Iterative Solver: Jacobi Method in OpenMP

```
#pragma omp parallel shared(a,b,x,new_x,n) private(iter,i,j,sum)
{
    #pragma omp for
    for (i = 0; i < n; i++)
        x[i] = b[i];
    for (iter = 0; iter < limit; iter++)
    {
        #pragma omp for
        for (i = 0; i < n; i++)
        {
            sum = -a[i][i] * x[i];
            for (j = 0; j < n; j++)
                sum = sum + a[i][j] * x[j];
            new_x[i] = (b[i] - sum) / a[i][i];
        }
        #pragma omp for
        for (i = 0; i < n; i++)
            x[i] = new_x[i];
    }
}
```

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$



OpenMP Iterative Solver Checking for Convergence

```
#pragma omp parallel shared(a,b,x,new_x,n,notdone) ...
{
    ...
    for (iter = 0; iter < limit; iter++)
    {
        ...
        #pragma omp for reduce(||:notdone) private(sum,i,j)
        for (i = 0; i < n; i++)
        {
            sum = 0;
            for (j = 0; j < n; j++)
                sum = sum + a[i][j] * x[j];
            if (fabs(sum - b[i]) >= tolerance)
                notdone = 1;
        }
        if (notdone == 0) break;
    }
}
```

*Can we change this to not
recomputing the sum?*

*Is there another solution
without reduce?*

$$\left| \sum_{j=0}^{n-1} a_{i,j} x_j^k - b_i \right| < \epsilon \quad \forall i = 0, \dots, n-1$$

Bertsekas and Tsitsiklis



OpenMP Iterative Solver Gauss-Seidel Relaxation

```
#pragma omp parallel shared(a,b,x,n,nt) private(iter,i,j,sum,tid,work,lo,hi,loc_x)
{
    #pragma omp single
        nt = omp_get_num_threads();
        tid = omp_get_thread_num();
        work = (n + nt-1) / nt;
        lo = work * tid;
        hi = lo + work;
        if (hi > n)
            hi = n;
        for (i = lo; i < hi; i++)
            x[i] = b[i];
        #pragma omp flush(x) // we need this?

        for (iter = 0; iter < limit; iter++)
        {
            #pragma omp barrier
            for (i = lo; i < hi; i++)
            {
                sum = -a[i][i] * x[i];
                for (j = 0; j < n; j++)
                {
                    if (j >= lo && j < i)
                        sum = sum + a[i][j] * loc_x[j-lo];
                    else
                        sum = sum + a[i][j] * x[j];
                }
                loc_x[i-lo] = (b[i] - sum) / a[i][i];
            }
            #pragma omp barrier
            for (i = lo; i < hi; i++)
                x[i] = loc_x[i-lo];
            #pragma omp flush(x) // we need this?
        }
    }
```

*Departure from pure
dataparallel model!*



Synchronous Computing with High-Performance Fortran

- High Performance Fortran (HPF) is an extension of Fortran 90 with constructs for parallel computing
 - Dataparallel FORALL
 - PURE (side-effect free functions)
 - Directives for recommended data distributions over processors
 - Library routines for parallel sum, prefix (scan), scattering, sorting, ...
- Uses the array syntax of Fortran 90 for as a dataparallel model of computation
 - Spreads the work of a single array computation over multiple processors
 - Allows efficient implementation on both SIMD and MIMD style architectures, shared memory and DSM
- But most users and vendors prefer OpenMP over HPF




HPF

```
!HPF$ PROCESSORS procname(dim1,...,dimN)  
!HPF$ DISTRIBUTE array1(dist),...,arrayM(dist) ONTO procname
```


Shaded area: P0 has X(1:2,1:8)

Example:


```
!HPF$ PROCESSORS pr(2,2)  
REAL X(8,8)  
!HPF$ DISTRIBUTE X(BLOCK,*) ONTO pr
```




(BLOCK, *)




(* , BLOCK)




(BLOCK, BLOCK)



(CYCLIC, *)



(CYCLIC, CYCLIC)









(CYCLIC, BLOCK)



HPF

!HPF\$ ALIGN array WITH target



Example:

REAL A(6),B(8),C(4)

REAL D(6,6),E(6,6)

!HPF\$ ALIGN A(I) WITH B(I)

Aligns arrays to target arrays to conform to target's data distribution over processors



HPF Heat Distribution Problem

$$h_{i,j} = \frac{h_{i-1,j} + h_{i+1,j} + h_{i,j-1} + h_{i,j+1}}{4}$$

```
!HPF$ PROCESSORS pr(4)
REAL h(100,100)
!HPF$ DISTRIBUTE h(BLOCK,*) ONTO pr
...
h(2:99,2:99) = 0.25*( h(1:98,2:99)+h(3:100,2:99)
                      +h(2:99,1:98)+h(2:99,3:100) )
```

Alternatively, with **FORALL**

```
FORALL (i=2:99,j=2:99) h(i,j) = 0.25*( h(i-1,j)+h(i+1,j)
                                         +h(i,j-1)+h(i,j+1) )
```

Remember: forall has “copy-in-copy-out” semantics, also HPF array assignments



HPF Heat Distribution Problem

Red-black Ordering

```
!HPF$ PROCESSORS pr(4)
      REAL h(100,100)
!HPF$ DISTRIBUTE h(BLOCK,*) ONTO pr
...
FORALL (i=2:99, j=2:99, MOD(i+j,2).EQ.0)
    h(i,j) = 0.25*(h(i-1,j)+h(i+1,j)+h(i,j-1)+h(i,j+1))
FORALL (i=2:99, j=2:99, MOD(i+j,2).EQ.1)
    h(i,j) = 0.25*(h(i-1,j)+h(i+1,j)+h(i,j-1)+h(i,j+1))
```




HPF Iterative Solver: Jacobi Method


```
!HPF$ PROCESSORS pr(4)
      REAL a(100,100),b(100),x(100)
!HPF$ ALIGN x(:) WITH a(:,*)
!HPF$ ALIGN b(:) WITH x(:)
!HPF$ ALIGN s(:) WITH x(:)
!HPF$ DISTRIBUTE a(BLOCK,*) ONTO pr
      x = b
      FORALL (i = 1:n) s(i) = SUM(a(i,:)*x(:))
      DO iter = 0,limit
        FORALL (i = 1:n) x(i) = (b(i) - s(i) + a(i,i)*x(i)) / a(i,i)
        FORALL (i = 1:n) s(i) = SUM(a(i,:)*x(:))
        IF (MAXVAL(ABS(s - b)) < tolerance) EXIT
      ENDDO
```

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

$$\left| \sum_{j=0}^{n-1} a_{i,j} x_j^k - b_i \right| < \epsilon \quad \forall i = 0, \dots, n-1$$



Gaussian Elimination



- The original system of equations is reduced to an upper triangular form
$$Ux = y$$
where U is a matrix of size $N \times N$ in which all elements below the diagonal are zero, and diagonal elements have the value 1
- Back substitution (Gauss-Jordan elimination): the new system of equations is solved to obtain the values of x

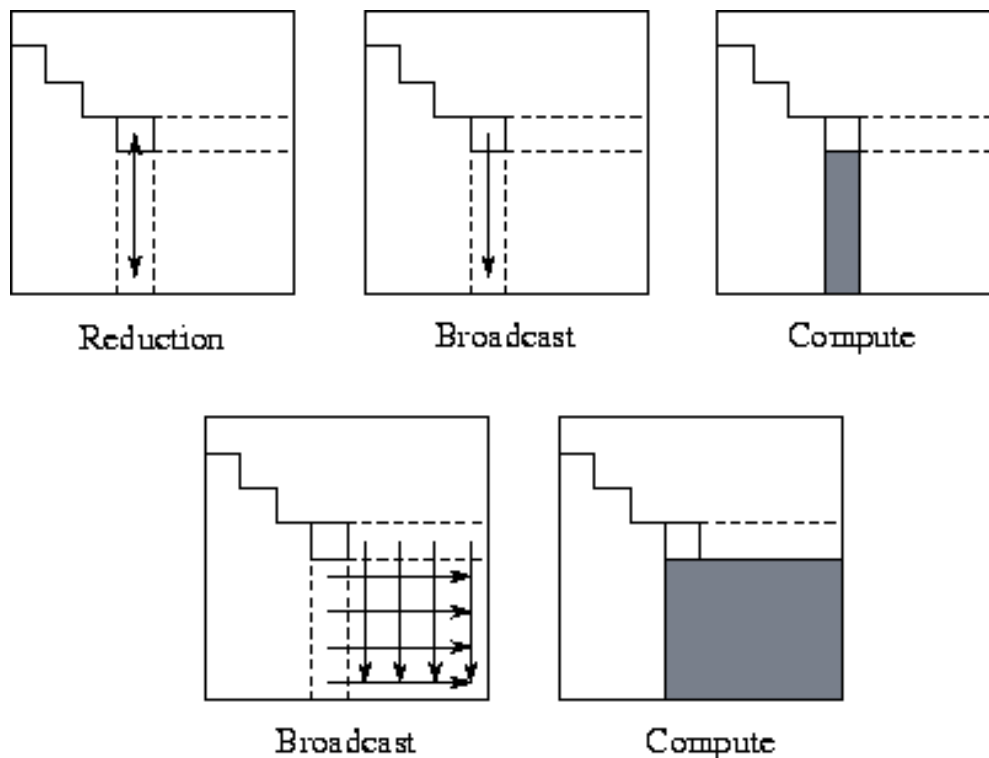
See: <http://www-unix.mcs.anl.gov/dbpp/text/node82.html>



HPF Gaussian Elimination 1

```
REAL A(n,n+1), X(n), Fac(n), Row(n+1)
INTEGER indx(n), itmp(1), max_indx, i, j, k
!HPF$ ALIGN Row(j) WITH A(1,j)
!HPF$ ALIGN X(i) WITH A(i,N+1)
!HPF$ DISTRIBUTE A(*,CYCLIC)
indx = 0
DO i = 1,n
    itmp = MAXLOC(ABS(A(:,i)), MASK=indx.EQ.0)      ! Stage 1
    max_indx = itmp(1)                               ! Stage 2
    indx(max_indx) = i
    Fac = A(:,i) / A(max_indx,i)                    ! Stage 3+4
    Row = A(max_indx,:)
    FORALL (j=1:n, k=i:n+1, indx(j).EQ.0)           ! Stage 5
        A(j,k) = A(j,k) - Fac(j)*Row(k)
ENDDO
! Row exchange
FORALL (j=1:n) A(indx(j),:) = A(j,:)
! Backsubstitution, uses B(:) stored in A(1:n,n+1)
DO j = n,1,-1
    X(j) = A(j,n+1) / A(j,j)
    A(1:j-1,n+1) = A(1:j-1,n+1) - A(1:j-1,j)*X(j)
ENDDO
```

HPF Gaussian Elimination 2



- Computing the upper triangular form takes five stages:
 1. Reduction with **MAXLOC**
 2. Broadcast (copy) **max_indx**
 3. Compute scale factors **Fac**
 4. Broadcast scale factor **Fac** and pivot row value **Row(k)**
 5. Row update with **FORALL**