

# Programming with Message Passing **PART I: Basics**

**HPC Fall 2012**

*Prof. Robert van Engelen*





# Overview

- Communicating processes
- MPMD and SPMD
- Point-to-point communications
  - Send and receive
  - Synchronous, blocking, and nonblocking message passing
  - Message selection
- Collective communications
  - broadcast, gather, scatter, barrier
- Further reading

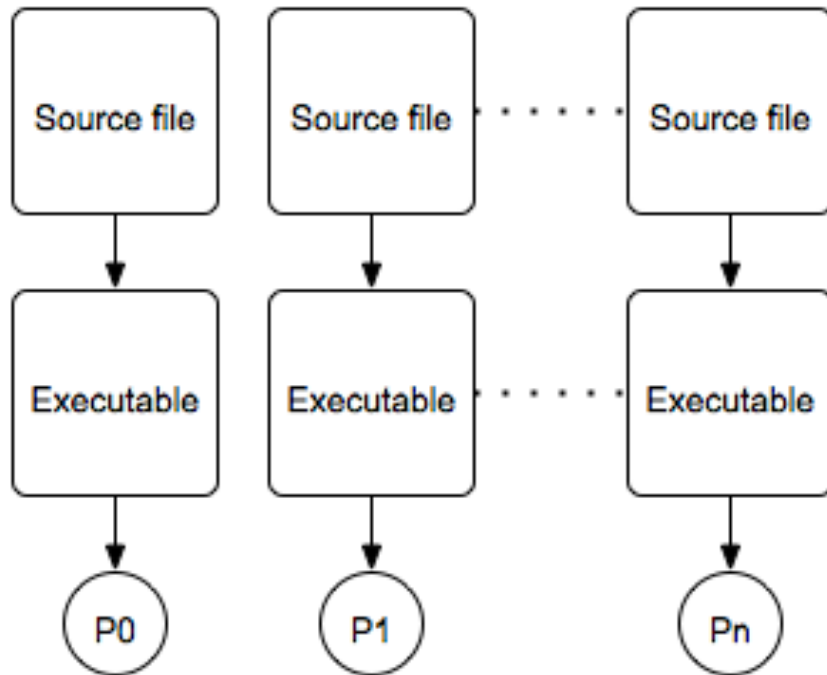


# Process Creation

- Processes communicate via message passing
- How are processes created?
  - *Static process creation*
    - All processes are specified before execution
    - Fixed number of processes executed
    - Example: `mpirun` command to start MPI program on  $n$  processors:  
`mpirun -np n`
  - *Dynamic process creation*
    - Processes are created during the execution of other processes
    - Processes can fork new processes
    - Management (start/stop), synchronization, and communication are more difficult

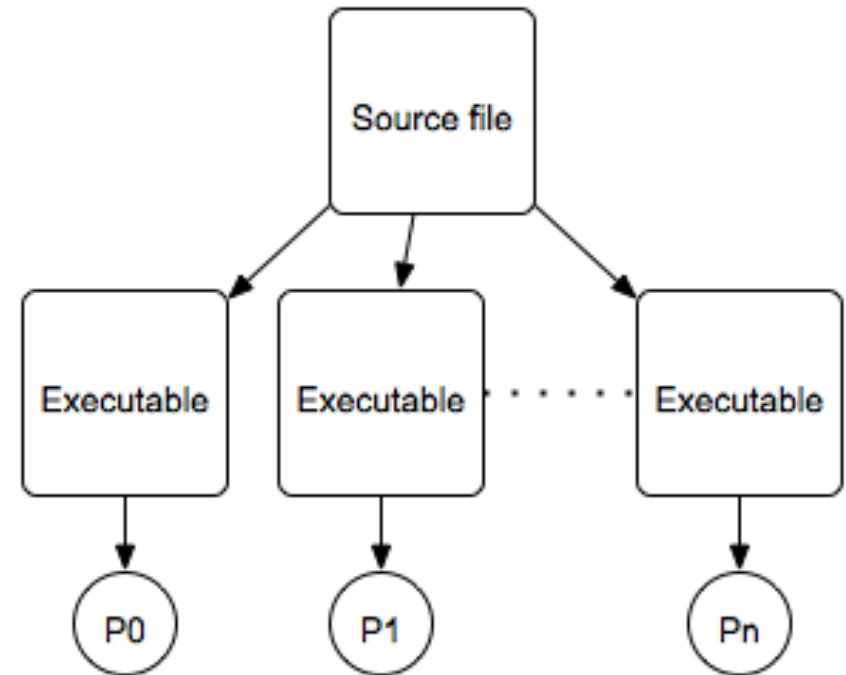
# MPMD Versus SPMD

*Multiple Program  
Multiple Data  
(MPMD)*



*Example: web server and web browsers*

*Single Program  
Multiple Data  
(SPMD)*



*Example: MPI program*



# Basic Send and Receive

- Send and receive operations w/o source and destination process ID
  - send(&x)**                      send **x** to any destination
  - recv(&y)**                      receive **y** from any source
- Send and receive operations with source and destination process ID
  - send(&x, destID)**                      send **x** to destination **destID**
  - recv(&y, srcID)**                      receive **y** from source **srcID**
- Data type of **x** and **y** must match
- What about *rendezvous*?
  - ☐ Should the sender wait until message is received by destination?



# Synchronous and (non)Blocking Send Operations

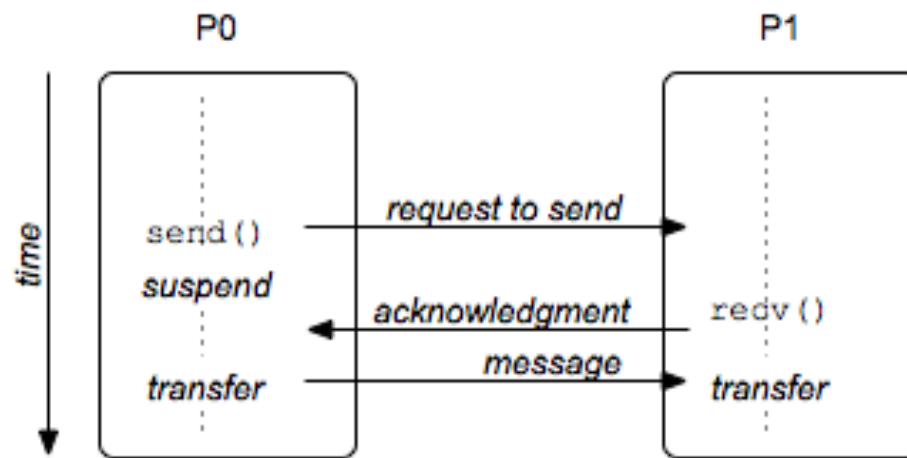
- *Synchronous (also called blocking)*
  - Both sender and receiver wait until entire message is delivered
- *(Locally) blocking send*
  - Sender sends x and may continue operating on x
  - Copy of x is buffered (causing process to be temporarily suspended until copy is completed) or immediately transmitted (when x is small)
  - A receiver may accept message at any time
- *Nonblocking send*
  - Sender initiates a “send” of x and immediately continues
  - Sender cannot further operate on x (data x is in *transfer state*)
  - Receiver may accept message at any time



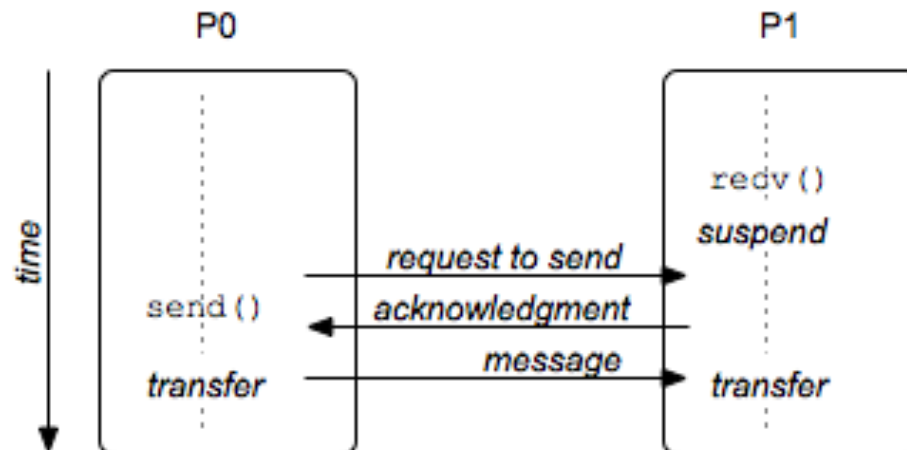
# Blocking and Nonblocking Receive Operations

- *Blocking receive*
  - Receiver waits for data to be completely transferred
  
- *Nonblocking receive*
  - Receiver indicates it is ready to receive data into y
  - A *handle* is returned that allows the receiver to query the status of the received data for y
  
- Note: any type of send can be paired with any type of receive

# Synchronous Send and Recv



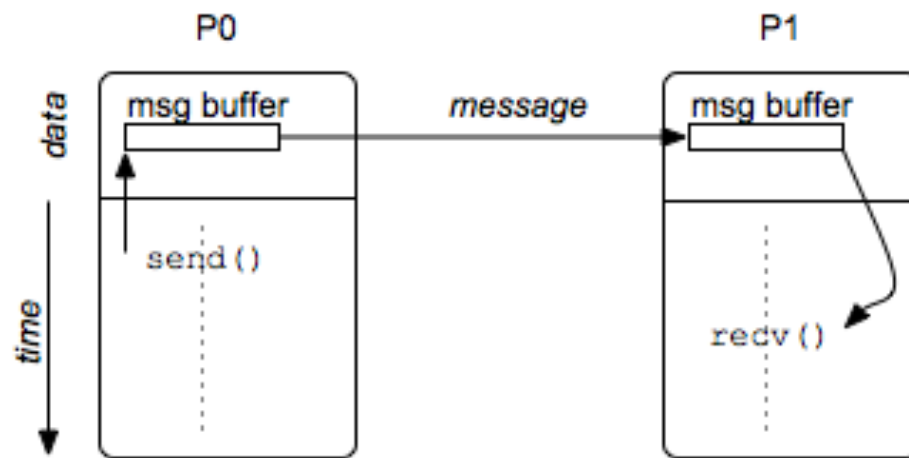
**send ()** occurs before **recv ()**  
*P0 is suspended until a receiver is ready*



**recv ()** occurs before **send ()**  
*P1 is suspended until a sender is ready*

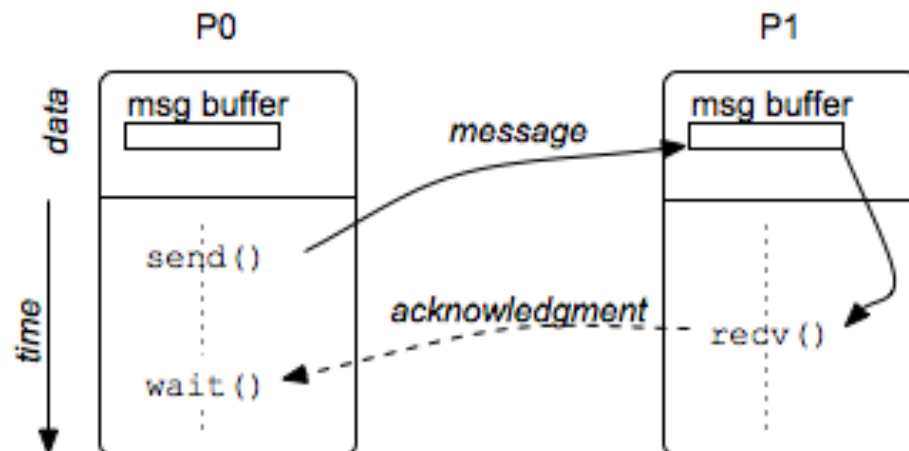


# (non)Blocking Send and Recv



*In a (locally) blocking **send()**, process P0 continues after the message is locally buffered or in transit to receiver, and it is safe for P0 to modify the data*


*Blocking: P0 suspends until a **recv()** is posted*



*In a nonblocking **send()**, process P0 immediately continues and executes while message is delivered (hides the messaging latency)*

*P0 cannot modify data in transit, explicitly probe message status or wait until message was received*

# Deadlock



*Deadlock with synchronous blocking send operations: both processors wait for data to be send to a receiver that is not ready to accept the message*

*Note: nonblocking sends and sendrecv() operations (send-recv exchanges) are safe to use for this example*




# Message Selection

- Send and receive operations indicate source/destination process ID
  - Id can be a wildcard
- What if multiple messages are *asynchronously transmitted out-of-order* to a destination?
  - Messages may be queuing up and end up being transmitted or accepted in different order, as if they “crossed” in transit
  - Cannot rely on message ordering with blocking/nonblocking send, even when sends are initiated by one processes
  - *Message tags* are used to match send and receive operations
    - send(&x, destID, tag)**
    - recv(&y, srcID, tag)**
    - message is transferred when **tag** value matches

# Broadcast


- *Multicast*: a root process sends a message to a specific subset of processes
- *Broadcast* = *multicast* within a process group
- First a group must be formed and root process selected






# Scatter

- *Scatter*: a root process sends elements of an array  $a[0, \dots, n]$  to the enumerated processes  $P_i$ ,  $i=0, \dots, n$
- First a group must be formed and root process selected




# Gather

- *Gather*: a root process collects data from the enumerated processes  $P_i$ ,  $i=0,\dots,n$  and puts them into the elements of an array  $a[0,\dots,n]$
- First a group must be formed and root process selected




# Reduce

- *Reduce*: a root process collects data from the enumerated processes  $P_i$ ,  $i=0,\dots,n$  and reduces it to a single value
- First a group must be formed and root process selected



# AllGather and AllReduce

- *AllGather and AllReduce*: perform gather/reduce and broadcast result
- First a group must be formed and root process selected






# Barrier

- *Barrier*: synchronization point


*Example barrier based on an allReduce  
(typically more efficient implementations are used)*




# Processor Groups and Interconnect Topologies

- A *processor group* is a subset of all processors
  - Collective communications occur within a group
- A group (including the group of all processors) can be mapped to a *virtual topology*
  - When the virtual topology of a group is matched to a *physical interconnect topology* that is a close approximation of the virtual topology, message latencies are more predictable


*Group 1 with 1D Cartesian virtual topology*



*Group 2 with 2D Cartesian virtual topology*



*interconnect topology*





# Further Reading

- [PP2] pages 42-51